

Wormhole: A Fast Ordered Index Data Structure

Xingbo Wu



Fan Ni Song Jiang




Index for In-memory Large-scale Data Stores

- The index can be very big.
 - The memory becomes increasingly large.
 - The indexed data (key-value items) can be small.
 - Many billions of keys have to be indexed.
- The index operations can be very expensive.


“The TPC-H queries spend up to 94% (35% on average) and TPC-DS queries spend up to 77% (45% on average) of their execution time on indexing.”
Kocberber, et al. at IEEE/ACM MICRO’13
- The index needs to be sorted.
 - Hash table with $O(1)$ operations often isn’t a choice.
 - **But we will leverage it to accelerate search in a sorted index.**

B+-tree Can be too Expensive



- $O(\log N)$ search cost (N is the number of keys)
 - ~30 key-comparisons for one billion keys.
- A key comparison may induce multiple cache misses.
- An index search may end up with 50-100 memory accesses.

Prefix Tree Can be Time- and Space-inefficient



- $O(L)$ search cost (L is token count in a search key).
 - 100+ tokens in a key (e.g, URL)
- High space cost due to small node size.

The Dilemma and our Solution

Data Structure	Pros	Cons
B+ Tree	<ul style="list-style-type: none">* Space efficient (large leaf nodes)* Support of range operations	High lookup cost with a large N
Prefix Tree	Lookup cost not correlated with N	<ul style="list-style-type: none">* High lookup cost even with a moderate L* Space inefficiency
Hash Table	$O(1)$ lookup cost	No support of range operations

We will orchestrate **B+-tree, prefix tree, and hash table** in one index structure that:

- has $O(\log L)$ search cost;
- is memory-efficient;
- supports range search.


A Closer Look at B+-tree

The tree of internal nodes:


- $O(\log M)$ lookup cost
- Operations for maintaining tree balance

The list of leaf nodes:

- All keys in sorted order
- Fast range operations
- Constant lookup cost within each node



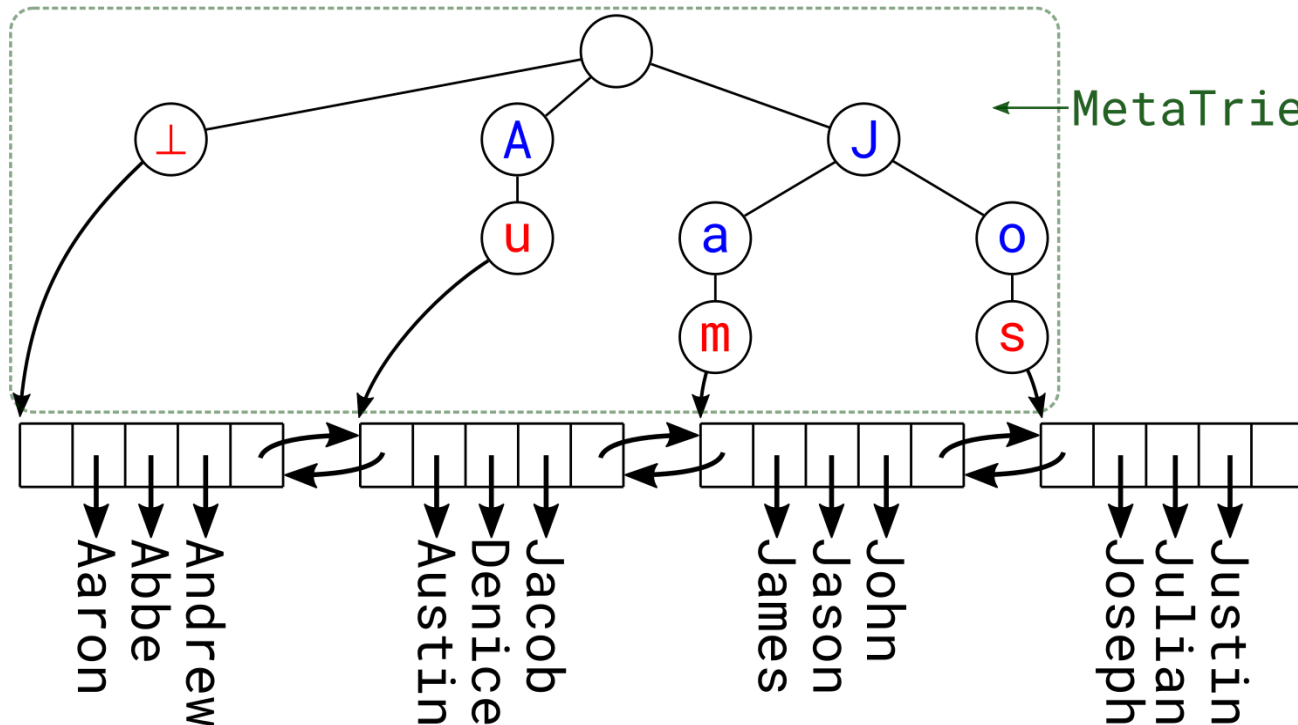
Replace MetaTree with a Hash Table?



However, this doesn't work:


- **No support of insert;**
- **No full support of range search.**

Replace MetaTree with MetaTrie (a Prefix Tree)




- Keys-at-its-left < **anchor** ≤ keys-at-its-right
 - Example anchors: “Au”, “Jam”, “Jos”
- All prefixes of anchors are in MetaTrie.


Search for “James”



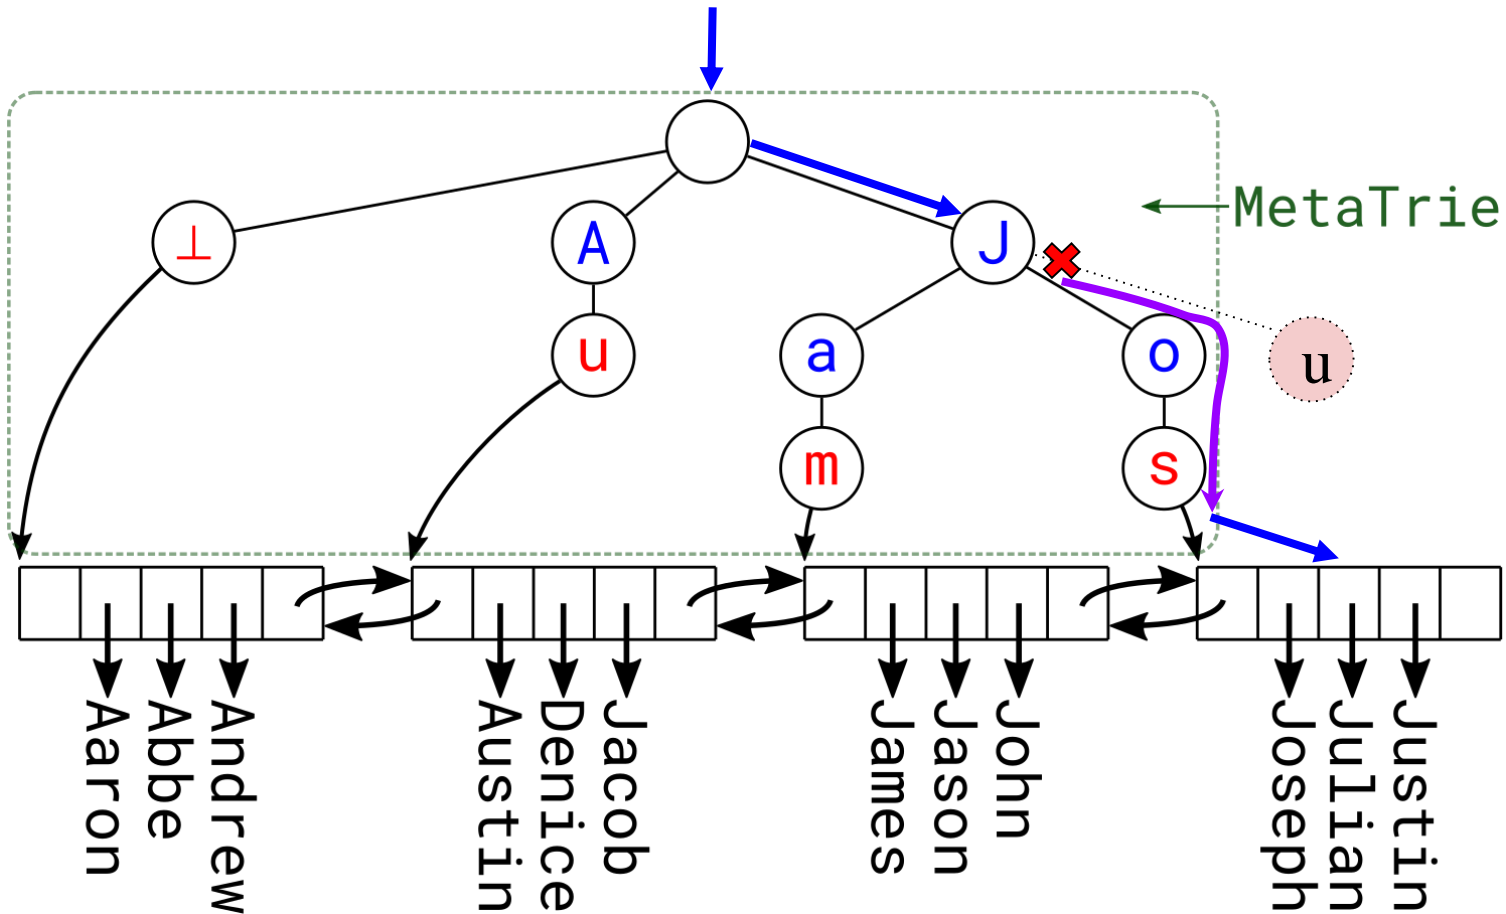
Search for “Denice”



Search for “Denice”




Search for “Julian”



But the Search Cost is Still $O(L)$

- Two phases in a search:
 - Find the longest prefix match (LPM) on the trie [1st phase].
 - Reach the right-most leaf of the missing node's left sibling, or the left-most leaf of the missing node's right sibling [2nd phase].

Binary search on the prefix length ($O(L) \rightarrow O(\log L)$)



Record pointers of left-most and right-most nodes at each trie node ($O(L) \rightarrow O(1)$)

Andrew
Abbe
Aaron

Jacob
Denice
Austin

John
Jason
James

Justin
Julian
Joseph

An Example: Binary Search on the Key with a Hash Table

- The anchor is “Alexander”
- All prefixes of the anchor are inserted to the hash table.
 - “A”, “Al”, “Ale”, ... , “Alexander”
- Binary search of LPM for key “Alexandria”
 - $\text{Len} = 5 \left(\lceil (0+9)/2 \rceil \right)$: “Alexa” exists.

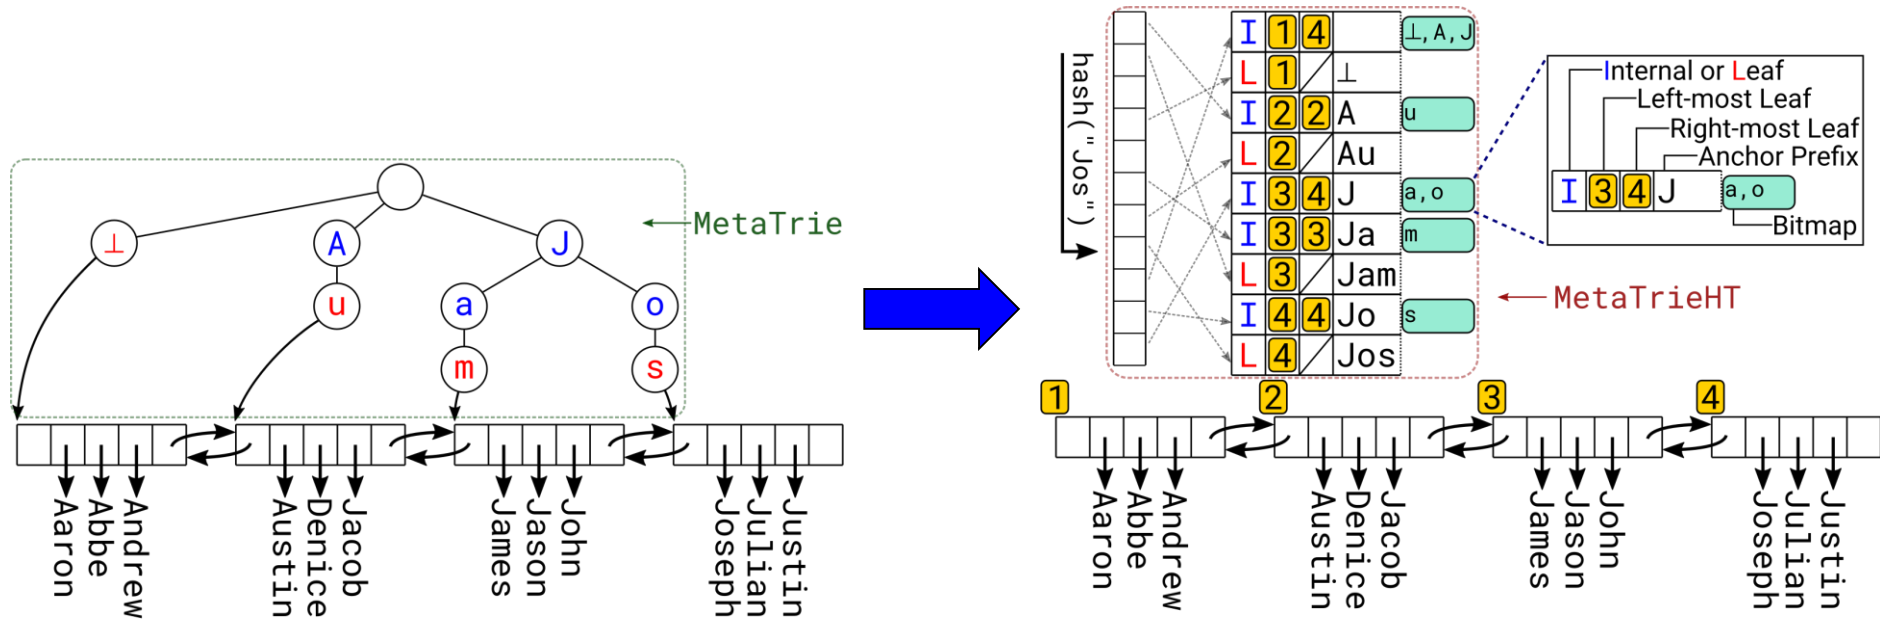
An Example: Binary Search on the Key with a Hash Table

- The anchor is “Alexander”
- All prefixes of the anchor are inserted to the hash table.
 - “A”, “Al”, “Ale”, ... , “Alexander”
- Binary search of LPM for key “Alexandria”
 - $\text{Len} = 5 \left(\lceil (0+9)/2 \rceil \right)$: “Alexa” exists.
 - $\text{Len} = 7 \left(\lceil (5+9)/2 \rceil \right)$: “Alexand” exists

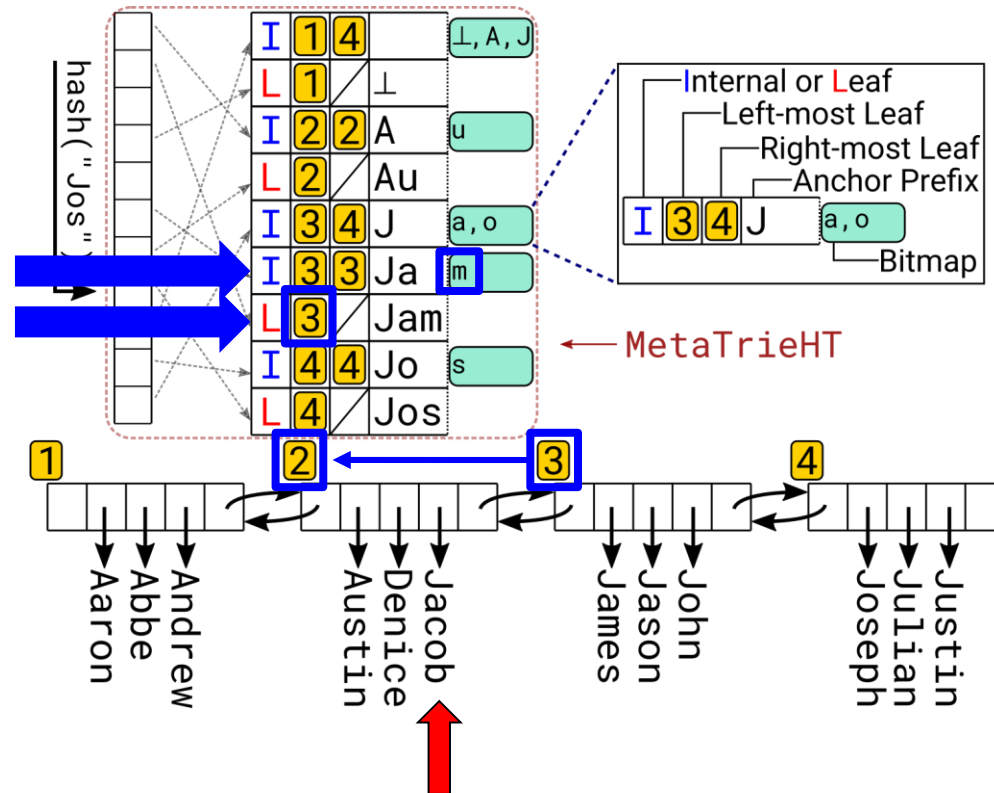
An Example: Binary Search on the Key with a Hash Table

- The anchor is “Alexander”
- All prefixes of the anchor are inserted to the hash table.
 - “A”, “Al”, “Ale”, ... , “Alexander”
- Binary search of LPM for key “Alexandra”
 - $\text{Len} = 5 \left(\left\lceil \frac{(0+9)}{2} \right\rceil \right)$: “Alexa” exists.
 - $\text{Len} = 7 \left(\left\lceil \frac{(5+9)}{2} \right\rceil \right)$: “Alexand” exists
 - $\text{Len} = 8 \left(\left\lceil \frac{(7+9)}{2} \right\rceil \right)$: “Alexandr” does not exist.
 - so LPM is “Alexand” of length 7.

Wormhole: Replace MetaTree with MetaTrieHT (a Prefix on a Hash Table)



Search for “Jacob”



- Search “**Ja**”
- But “**Jac**” is not in the table.
- Go to **Jam**’s left-most leaf node, which is Node 3
- Go to the node’s left sibling, which is Node 2
- Search for “**Jacob**” in Node 2.


A Recap of Wormhole

- Its lookup cost is $O(\log L)$
 - Asymptotically (much) better than $O(\log N)$ and $O(L)$
- Its space cost is comparable to B+ tree
 - One anchor per leaf node
- Efficient support of range Search
- Conduct split/Merge of leaf nodes (for insertion/deletion)
 - Same as B+-tree (But do not need to maintain tree balance)
 - Efficient concurrency control (see the paper)

Performance Evaluation


- The workload
 - Amazon reviews metadata (143M keys)
 - URLs from (“http://zwischenruf.at/?p=1012”) (192M keys)
 - Randomly generated keys of different sizes (8B, 16B, ... 1024B)
- The server
 - 16-core Intel Xeon E5-2697A, 40MB shared LLC, 256GB DRAM@2400MHz
- Indexes in comparison
 - Skiplist
 - B+-tree
 - Masstree
 - Cuckoo Hash Table

Lookup Performance



Compare to Cuckoo Hash Table (w/ Point Search)

ONLY 2x ~ 3x slower



Conclusions

- A new index data structure for sorted keys with an asymptotically low cost of **$O(\log L)$** .
- A well optimized implementation delivers **throughput multiple times higher**.
- A promising data structure for fine-grain-indexed big data store.

Source code is available for downloading at
[*https://github.com/wuxb45/wormhole*](https://github.com/wuxb45/wormhole)